

Churn Customer Prediction

September 13, 2025

1 Customer Churn Prediction

This notebook demonstrates how to build a **supervised machine learning model** to predict customer churn, an online food ordering provider.

Why churn prediction matters? - Acquiring a new customer is 5–7 times more expensive than retaining an existing one. - By predicting churn, can target discounts, loyalty campaigns, or personalized offers to customers **before they leave**.

1.1 Step 1: Import Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from sklearn.metrics import classification_report, confusion_matrix, \
    roc_auc_score, roc_curve
```

1.2 Step 2: Create Sample Dataset

```
[2]: # Simulated dataset
np.random.seed(42)
n = 1000

data = pd.DataFrame({
    'customer_id': range(1, n+1),
    'last_order_days_ago': np.random.randint(1, 365, n),
    'total_orders': np.random.poisson(10, n),
    'avg_order_value': np.random.uniform(5, 50, n),
    'complaints': np.random.binomial(5, 0.1, n),
    'discount_used': np.random.randint(0, 20, n),
```

```
'churn': np.random.binomial(1, 0.3, n) # 1 = churned, 0 = active
})

data.head()
```

```
[2]:  customer_id  last_order_days_ago  total_orders  avg_order_value  \
0           1           103           13      24.118288
1           2           349           11      45.852733
2           3           271           9       27.817684
3           4           107           18      13.456417
4           5            72            7       8.463725

      complaints  discount_used  churn
0             1             6      0
1             0            17      0
2             2             6      0
3             0            16      1
4             1             6      0
```

1.3 Step 3: Exploratory Data Analysis

```
[3]: # Basic statistics
print(data.describe())


# Churn distribution
sns.countplot(x='churn', data=data)
plt.title("Churn Distribution")
plt.show()


# Correlation heatmap
plt.figure(figsize=(8,6))
sns.heatmap(data.drop('customer_id', axis=1).corr(), annot=True, cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()
```

```
      customer_id  last_order_days_ago  total_orders  avg_order_value  \
count  1000.000000      1000.000000      1000.000000      1000.000000
mean    500.500000      181.374000       10.064000       27.272050
std    288.819436      103.360018        3.204399       12.945629
min       1.000000        1.000000        2.000000        5.001382
25%     250.750000        97.750000        8.000000       16.227675
50%     500.500000       180.000000       10.000000       27.098596
75%     750.250000       268.000000       12.000000       38.286977
max    1000.000000       364.000000       22.000000       49.980097

      complaints  discount_used      churn
count  1000.000000      1000.000000  1000.000000
mean     0.484000        9.375000    0.296000
```

std	0.675426	5.827034	0.456719
min	0.000000	0.000000	0.000000
25%	0.000000	4.000000	0.000000
50%	0.000000	9.000000	0.000000
75%	1.000000	15.000000	1.000000
max	3.000000	19.000000	1.000000





1.4 Step 4: Data Preprocessing

```
[4]: X = data.drop(['customer_id', 'churn'], axis=1)
y = data['churn']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42, stratify=y)

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

1.5 Step 5: Train Models

```
[5]: # Logistic Regression
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train_scaled, y_train)
y_pred_log = log_reg.predict(X_test_scaled)

# Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

# XGBoost
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
    ↪random_state=42)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)
```

```
C:\Users\SPINO SHOP\anaconda3\Lib\site-packages\xgboost\training.py:183:
UserWarning: [12:44:55] WARNING: C:\actions-
runner\work\xgboost\xgboost\src\learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

1.6 Step 6: Model Evaluation

```
[6]: print("Logistic Regression:\n", classification_report(y_test, y_pred_log))
print("Random Forest:\n", classification_report(y_test, y_pred_rf))
print("XGBoost:\n", classification_report(y_test, y_pred_xgb))

# Confusion Matrix for Random Forest
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d',
    ↪cmap='Blues')
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# ROC Curve for XGBoost
y_probs = xgb.predict_proba(X_test)[: ,1]
fpr, tpr, _ = roc_curve(y_test, y_probs)
plt.plot(fpr, tpr, label=f"AUC = {roc_auc_score(y_test, y_probs):.2f}")
plt.plot([0,1], [0,1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - XGBoost")
plt.legend()
```

```
plt.show()
```

```
C:\Users\SPINO SHOP\anaconda3\Lib\site-  
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
C:\Users\SPINO SHOP\anaconda3\Lib\site-  
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
C:\Users\SPINO SHOP\anaconda3\Lib\site-  
packages\sklearn\metrics\_classification.py:1565: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Logistic Regression:


	precision	recall	f1-score	support
0	0.70	1.00	0.83	141
1	0.00	0.00	0.00	59
accuracy			0.70	200
macro avg	0.35	0.50	0.41	200
weighted avg	0.50	0.70	0.58	200


Random Forest:

	precision	recall	f1-score	support
0	0.69	0.91	0.79	141
1	0.13	0.03	0.05	59
accuracy			0.65	200
macro avg	0.41	0.47	0.42	200
weighted avg	0.53	0.65	0.57	200

XGBoost:


	precision	recall	f1-score	support
0	0.69	0.75	0.72	141
1	0.26	0.20	0.23	59
accuracy			0.59	200
macro avg	0.47	0.48	0.47	200
weighted avg	0.56	0.59	0.58	200





1.7 Step 7: Feature Importance

```
[7]: importances = rf.feature_importances_  
     feat_names = X.columns  
  
     feat_imp = pd.DataFrame({'Feature': feat_names, 'Importance': importances})  
     feat_imp.sort_values(by='Importance', ascending=False, inplace=True)  
  
     sns.barplot(x='Importance', y='Feature', data=feat_imp)  
     plt.title("Random Forest Feature Importance")  
     plt.show()
```

1.8 Step 8: Conclusion & Business Use Case

- The models give us a way to predict which customers are likely to churn.
- We can integrate this model into the CRM/marketing system:
 - Send targeted discounts to **high-risk churners**.
 - Recommend personalized meals to re-engage customers.
 - Monitor churn rates across restaurants and regions.

Next Steps: - Replace synthetic dataset with real customer order data.

- Regularly retrain the model with new data.
- Deploy as an API to integrate with the app.