

# Churn Customer Prediction

September 16, 2025

## 1 Customer Churn Prediction

This notebook demonstrates how to build a **supervised machine learning model** to predict customer churn, an online food ordering provider.

**Why churn prediction matters?** - Acquiring a new customer is 5–7 times more expensive than retaining an existing one. - By predicting churn, can target discounts, loyalty campaigns, or personalized offers to customers **before they leave**.

### 1.1 Step 1: Import Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from sklearn.metrics import classification_report, confusion_matrix, \
    roc_auc_score, roc_curve

import warnings
warnings.filterwarnings('ignore')
```

### 1.2 Step 2: Create Sample Dataset

```
[2]: # Simulated dataset
np.random.seed(42)
n = 10000

data = pd.DataFrame({
    'customer_id': range(1, n+1),
    'last_order_days_ago': np.random.randint(1, 365, n),
    'total_orders': np.random.poisson(10, n),
```

```

    'avg_order_value': np.random.uniform(5, 50, n),
    'complaints': np.random.binomial(5, 0.1, n),
    'discount_used': np.random.randint(0, 20, n),
    'churn': np.random.binomial(1, 0.3, n) # 1 = churned, 0 = active
})

data.head()

```

```

[2]:
customer_id  last_order_days_ago  total_orders  avg_order_value  \
0           1                  103             10      38.251680
1           2                  349             9       37.210404
2           3                  271             9       36.023084
3           4                  107             9       17.451155
4           5                   72             9       30.801515

complaints  discount_used  churn
0           1              4      1
1           0             13      1
2           0              7      1
3           0             10      0
4           1             14      1

```

### 1.3 Step 3: Exploratory Data Analysis

```

[3]: # Basic statistics
print(data.describe())

# Churn distribution
sns.countplot(x='churn', data=data)
plt.title("Churn Distribution")
plt.show()

# Correlation heatmap
plt.figure(figsize=(8,6))
sns.heatmap(data.drop('customer_id', axis=1).corr(), annot=True,
            cmap='coolwarm')
plt.title("Feature Correlation")
plt.show()


```

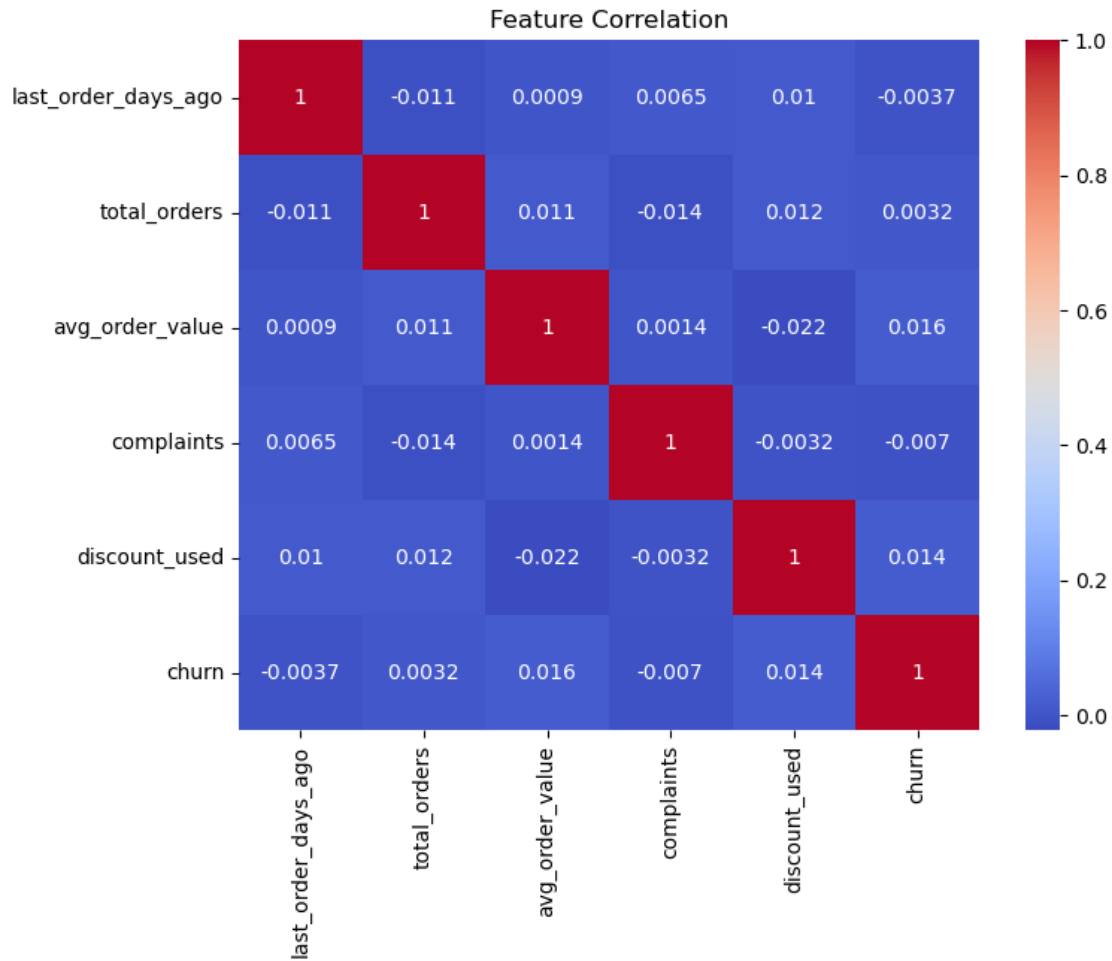
```

count  customer_id  last_order_days_ago  total_orders  avg_order_value  \
mean    5000.50000      181.922300      10.042600      27.451422
std     2886.89568      104.853453       3.160472      13.000012
min       1.00000       1.000000       0.000000       5.009142
25%     2500.75000       93.000000       8.000000      16.280541
50%     5000.50000      181.000000      10.000000      27.375755
75%     7500.25000      273.000000      12.000000      38.774759
max     10000.00000     364.000000      23.000000      49.995486

```

	complaints	discount_used	churn
count	10000.000000	10000.000000	10000.000000
mean	0.499400	9.437900	0.305100
std	0.668763	5.735227	0.460473
min	0.000000	0.000000	0.000000
25%	0.000000	4.000000	0.000000
50%	0.000000	9.000000	0.000000
75%	1.000000	14.000000	1.000000
max	4.000000	19.000000	1.000000





#### 1.4 Step 4: Data Preprocessing

```
[4]: X = data.drop(['customer_id', 'churn'], axis=1)
y = data['churn']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42, stratify=y)

# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## 1.5 Step 5: Train Models

```
[5]: # Logistic Regression
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train_scaled, y_train)
y_pred_log = log_reg.predict(X_test_scaled)

# Random Forest
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

# XGBoost
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
    ↪random_state=42)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)
```

## 1.6 Step 6: Model Evaluation

```
[6]: print("Logistic Regression:\n", classification_report(y_test, y_pred_log))
print("Random Forest:\n", classification_report(y_test, y_pred_rf))
print("XGBoost:\n", classification_report(y_test, y_pred_xgb))

# Confusion Matrix for Random Forest
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d',
    ↪cmap='Blues')
plt.title("Confusion Matrix - Random Forest")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# ROC Curve for XGBoost
y_probs = xgb.predict_proba(X_test)[:,-1]
fpr, tpr, _ = roc_curve(y_test, y_probs)
plt.plot(fpr, tpr, label=f"AUC = {roc_auc_score(y_test, y_probs):.2f}")
plt.plot([0,1], [0,1], 'k--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - XGBoost")
plt.legend()
plt.show()
```

Logistic Regression:

	precision	recall	f1-score	support
0	0.69	1.00	0.82	1390
1	0.00	0.00	0.00	610

accuracy			0.69	2000
macro avg	0.35	0.50	0.41	2000
weighted avg	0.48	0.69	0.57	2000

Random Forest:

	precision	recall	f1-score	support
0	0.70	0.92	0.79	1390
1	0.31	0.09	0.13	610


accuracy			0.66	2000
macro avg	0.50	0.50	0.46	2000
weighted avg	0.58	0.66	0.59	2000


XGBoost:

	precision	recall	f1-score	support
0	0.70	0.89	0.78	1390
1	0.32	0.12	0.17	610

accuracy			0.66	2000
macro avg	0.51	0.50	0.48	2000
weighted avg	0.58	0.66	0.60	2000





```
[7]: ex_data = [186, 8, 25.8733, 2, 11]

# Convert to numpy array and reshape to (1, n_features)
ex_data = np.array(ex_data).reshape(1, -1)

# Preprocessing by Standard Scaler
ex_data_scaled = scaler.transform(ex_data)

# Prediction By Algorithms:
ex_data_pred_log = log_reg.predict(ex_data_scaled)
ex_data_pred_rf = rf.predict(ex_data_scaled)
ex_data_pred_xgb = xgb.predict(ex_data_scaled)


print(f'Prediction by Logistic Regression is {ex_data_pred_log[0]}')
print(f'Prediction by Random Forest is {ex_data_pred_rf[0]}')
print(f'Prediction by XGBoost is {ex_data_pred_xgb[0]}')
```

```
Prediction by Logistic Regression is 0
Prediction by Random Forest is 1
Prediction by XGBoost is 0
```



## 1.7 Step 7: Feature Importance

```
[9]: importances = rf.feature_importances_  
     feat_names = X.columns  
  
     feat_imp = pd.DataFrame({'Feature': feat_names, 'Importance': importances})  
     feat_imp.sort_values(by='Importance', ascending=False, inplace=True)  
  
     sns.barplot(x='Importance', y='Feature', data=feat_imp)  
     plt.title("Logistic Regression Feature Importance")  
     plt.show()
```



## 1.8 Step 8: Conclusion & Business Use Case

- The models give us a way to predict which customers are likely to churn.
- We can integrate this model into the CRM/marketing system:
  - Send targeted discounts to **high-risk churners**.
  - Recommend personalized meals to re-engage customers.
  - Monitor churn rates across restaurants and regions.

**Next Steps:** - Replace synthetic dataset with real customer order data.  
- Regularly retrain the model with new data.

- Deploy as an API to integrate with the app.

[ ]: